


1. Project Title and Team Members
 - i. A Fused and IO-Efficient Attention Optimization for Transformer Inference on Llama3
 - ii. Team Members
 - i. Lyuyongkang Yuan (ly2188@nyu.edu)
 - ii. Junwei Yan (jy4831@nyu.edu)
2. Project Milestones (Breakdown of Major Steps)
 - i. Understand Llama 3 Attention Implementation
 - i. Study Meta's official Llama 3 attention code (ColumnParallelLinear, KV cache design, RoPE)
 - ii. Profiling Baseline Attention
 - i. Use PyTorch profiler and Nsight to capture kernel timeline for naïve attention
 - ii. Identify kernel patterns: Q/K/V GEMM -> QK -> scaling -> mask -> softmax -> P@V -> output GEMM
 - iii. Analyze Bottlenecks and Understand FlashAttention
 - i. Locate memory bottleneck and excessive global memory traffic from materializing attention scores
 - ii. Understand how FlashAttention reduce $O(T^2)$ IO and fuses softmax + matmul
 - iv. Implement a Simplified Fused Attention Kernel
 - i. Reproduce FlashAttention style tiling and online softmax in CUDA
 - ii. Test fused kernel in isolation (not integrated into Llama yet)
 - v. Model Level Integration ("Hot swapping" Llama Attention)
 - i. Replace Llama-3.2-1B's attention layer with custom fused attention module
 - ii. Measure token/s improvement
 - vi. Document the Roles of KV Cache and Quantization
 - i. Original Llama already support KV cache + mixed precision
 - ii. We will explain their latency benefits even if not re implemented from scratch
3. Milestones Completed and Main Results Obtained
 - i. Fully analyzed Llama 3 attention logic
 - i. We studied Meta's official repo and reproduced the attention flow in a simplified Python/CUDA prototype
 - ii. We validated the functional structure:
 $Q = WqX, K = WkX, V = WvX \rightarrow \text{reshape} \rightarrow \text{head partition} \rightarrow \text{repeat_kv} \rightarrow \text{attention softmax} \rightarrow \text{projection } Wo$

ii. Completed baseline profiling

i. Using PyTorch profiler

1. Each attention layer launches 15 kernels
2. Between kernels are micro gaps due to cuLaunchKernel dispatch overhead
3. Attention score matrix (T^2 heads) causes hundreds of MB of global memory reads and writes when seq_len is large
4. We generated a trace showing:
5. linear → linear → linear → unsqueeze → expand → transpose → matmul → div → add → softmax → matmul → transpose → contiguous → linear
6. One round attention calculation take ~40ms



4.

i. Identified bottlenecks

- i. Materializing the full attention matrix causes extreme memory IO (dominant cost at long sequence)
 - ii. Launching many small kernels leads to launch overhead accumulation
 - iii. HuggingFace uses scaled_dot_product_attention, while Meta's version uses multiple matmuls, both show the same IO issues
- Main insight so far:
The core bottleneck is IO-bound, not compute bound. Fusing scores + softmax + PV reduces $O(T^2)$ memory movement and eliminates multiple kernel launches.

5. Bottlenecks in Completing Remaining Milestones

i. Hard to analyze Meta's original attention in isolation

- i. Meta's attention uses: ColumnParallelLinear and RowParallelLinear; Distributed KV cache management; FSDP style tensor sharding assumptions
 - ii. Running the module standalone (without full model context) complicates profiling
 - iii. We addressed this by: Rebuilding a clean, de-parallelized reference attention module (equivalent to Meta's logic but purely single GPU, using regular Linear ops); Profiling this standalone version to get clean operator-level traces for differential comparison.
- ii. Integrating custom CUDA ops with PyTorch
 - i. Setting up C++/CUDA extension + autograd-safe forward path is non trivial
 - ii. Ensuring correct tensor shapes, dispatch types, and avoiding unnecessary Python overhead still requires time.
 - iii. Online softmax and tiling correctness
 - i. FlashAttention involves complex tile wise numerically stable softmax. Maintaining stability in FP16 while matching Llama's behavior is challenging.
- 6. Contribution of Each Team Member
 - i. Lyuyongkang Yuan
 - i. Write simplified attention reference module referring to original Llama 3 attention code
 - ii. Designed profiler experiments and produced PyTorch traces
 - iii. Analyzed baseline bottlenecks and mapped ops to attention phases
 - ii. Junwei Yan
 - i. Investigated FlashAttention kernel design
 - ii. Built profiling environment on local GPUs
 - iii. Evaluated memory traffic patterns and documented kernel launch overhead.
 - 7. Planned Work For Final Report
 - i. A working fused attention kernel (tile based, online softmax)
 - ii. End to end inference bench mark (Pytorch baseline vs fused kernel)
 - iii. Analyze of kernel timeline reduction and IO reduction
 - iv. Explanation of how KV cache and quantization complement attention fusion